

EduMIPS64

version 1.2.7.1

Andrea Spadaccini and the EduMIPS64 development team

April 13, 2020

Contents

Welcome to the EduMIPS64 documentation!	1
Source files format	1
Memory limits	1
The <i>.data</i> section	1
The <i>.code</i> section	2
The <i>#include</i> command	4
The instruction set	4
ALU Instructions	4
Load/Store instructions	7
Flow control instructions	8
The <i>SYSCALL</i> instruction	9
<i>SYSCALL 0 - exit()</i>	9
<i>SYSCALL 1 - open()</i>	9
<i>SYSCALL 2 - close()</i>	10
<i>SYSCALL 3 - read()</i>	10
<i>SYSCALL 4 - write()</i>	10
<i>SYSCALL 5 - printf()</i>	10
Other instructions	11
<i>BREAK</i>	11
<i>NOP</i>	11
<i>TRAP</i>	11
<i>HALT</i>	11
Floating Point Unit	11
Special values	11
NaN or Invalid Operation	11
Zeroes or Underflows	12
Infinites or Overflows	12
Exception configuration	12
The <i>.double</i> directive	12
The FCSR register	12
Instruction set	12
The user interface	17
The menu bar	17
File	17
Execute	17
Configure	18
Tools	18
Window	18
Help	18
Frames	18

Cycles	18
Registers	18
Statistics	18
Pipeline	18
Memory	19
Code	19
Input/Output	19
Dialogs	19
Settings	19
Dinero Frontend	19
Help	20
Command line options	20
Running EduMIPS64	20
Code Examples	20
SYSCALL	20
SYSCALL 0	20
SYSCALL 1	20
SYSCALL 2	21
SYSCALL 3	22
SYSCALL 4	22
SYSCALL 5	23
A more complex usage example of SYSCALL 5	23

Welcome to the EduMIPS64 documentation!

EduMIPS64 is a MIPS64 Instruction Set Architecture (ISA) simulator. It is designed to be used to execute small programs that use the subset of the MIPS64 ISA implemented by the simulator, allowing the user to see how instructions behave in the pipeline, how stalls are handled by the CPU, the status of registers and memory and much more. It is both a simulator and a visual debugger.

The website for the project is <http://www.edumips.org>, and the code is hosted at <http://github.com/lupino3/edumips64>. If you find any bugs, or have any suggestion for improving the simulator, please open an issue on github or send an email at bugs@edumips.org.

EduMIPS64 is developed by a group of students of the University of Catania (Italy), and started as a clone of WinMIPS64, even if now there are lots of differences between the two simulators.

This manual will introduce you to EduMIPS64, and will cover some details on how to use it.

The first chapter of this manual covers the format of source files accepted by the simulator, describing the data types and the directives, in addition to command line parameters. In the second chapter there's an overview of the subset of the MIPS64 instruction set that is accepted by EduMIPS64, with all the needed parameters and indications to use them. The third chapter is a description of the Floating Point Unit and its instruction set. The fourth chapter is a description of the user interface of EduMIPS64, that explains the purpose of each frame and menu, along with a description of the configuration dialog, the Dinero frontend dialog, the Manual dialog and command line options. The fifth chapter contains some useful examples.

This manual describes EduMIPS64 version 1.2.7.1.

Source files format

EduMIPS64 tries to follow the conventions used in other MIPS64 and DLX simulators, so that old time users will not be confused by its syntax.

There are two sections in a source file, the *data* section and the *code* section, introduced respectively by the *.data* and the *.code* directives. In the following listing you can see a very basic EduMIPS64 program:

```
; This is a comment
.data
label: .word 15      ; This is an inline comment

.code
daddi r1, r0, 0
syscall 0
```

To distinguish the various parts of each source code line, any combination of spaces and tabs can be used, as the parser ignores multiple spaces and only detects whitespaces to separate tokens.

Comments can be specified using the “;” character, everything that follows that character will be ignored. So a comment can be used “inline” (after the directive) or on a row by itself.

Labels can be used in the code to reference a memory cell or an instruction. They are case insensitive. Only a label for each source code line can be used. The label can be specified one or more rows above the effective data declaration or instruction, provided that there's nothing, except for comments and empty lines, between the label and the declaration.

Memory limits

EduMIPS64 has a fixed memory size for both data (the *.data* section, capped at 640 kB – i.e., 80000 64-bit values) and instructions (the *.code* section, capped at 128 kB – i.e., 32000 instructions, each occupying 32 bits).

These limits are hardcoded in the simulator.

The *.data* section

The *data* section contains commands that specify how the memory must be filled before program execution starts. The general form of a *.data* command is:

```
[label:] .datatype value1 [, value2 [, ...]]
```

EduMIPS64 supports different data types, that are described in the following table.

Type	Directive	Bits required
Byte	<i>.byte</i>	8
Half word	<i>.word16</i>	16
Word	<i>.word32</i>	32
Double Word	<i>.word</i> or <i>.word64</i>	64

Please note that a double word can be introduced either by the *.word* directive or by the *.word64* directive.

All the data types are interpreted as signed. This means that integer literals in the *.data* section must be between $-2^{(n-1)}$ and $2^{(n-1)} - 1$ (inclusive).

There is a big difference between declaring a list of data elements using a single directive or by using multiple directives of the same type. EduMIPS64 starts writing from the next 64-bit double word as soon as it finds a datatype identifier, so the first *.byte* statement in the following listing will put the numbers 1, 2, 3 and 4 in the space of 4 bytes, taking 32 bits, while code in the next four rows will put each number in a different memory cell, occupying 32 bytes:

```
.data
.byte    1, 2, 3, 4
.byte    1
.byte    2
.byte    3
.byte    4
```

In the following table, the memory is represented using byte-sized cells and each row is 64 bits wide. The address on the left side of each row of the table refers to the right-most memory cell, that has the lowest address of the eight cells in each line.

0	0	0	0	0	4	3	2	1
8	0	0	0	0	0	0	0	1
16	0	0	0	0	0	0	0	2
24	0	0	0	0	0	0	0	3
36	0	0	0	0	0	0	0	4

There are some special directives that need to be discussed: *.space*, *.ascii* and *.asciiz*.

The *.space* directive is used to leave some free space in memory. It accepts as a parameter an integer, that indicates the number of bytes that must be left empty. It is handy when you must save some space in memory for the results of your computations.

The *.ascii* directive accepts strings containing any of the ASCII characters, and some special C-like escaping sequences, that are described in the following table, and puts those strings in memory.

Escaping sequence	Meaning	ASCII code
\0	Null byte	0
\t	Horizontal tabulation	9
\n	Newline character	10
"	Literal quote character	34
\	Literal backslash character	92

The *.asciiz* directive behaves exactly like the *.ascii* command, with the difference that it automatically ends the string with a null byte.

The .code section

The *code* section contains commands that specify how the memory must be filled when the program will start. The general form of a *.code* command is:

```
[label:] instruction [param1 [, param2 [, param3]]]
```

The *code* section can be specified with the *.text* alias.

The number and the type of parameters depends on the instruction itself.

Instructions can take three types of parameters:

- *Registers* a register parameter is indicated by an uppercase or lowercase “r”, or a \$, followed by the number of the register (between 0 and 31), as in “r4”, “R4” or “\$4”;
- *Immediate values* an immediate value can be a number or a label; the number can be specified in base 10 or in base 16: base 10 numbers are simply inserted by writing the number, while base 16 number are inserted by putting before the number the prefix “0x”
- *Address* an address is composed by an immediate value followed by a register name enclosed in brackets. The value of the register will be used as base, the value of the immediate will be the offset.

The size of immediate values is limited by the number of bits that are available in the bit encoding of the instruction.

You can use standard MIPS assembly aliases to address the first 32 registers, appending the alias to one of the standard register prefixes like “r”, “\$” and “R”. See the next table.

Register	Alias
0	<i>zero</i>
1	<i>at</i>
2	<i>v0</i>
3	<i>v1</i>
4	<i>a0</i>
5	<i>a1</i>
6	<i>a2</i>
7	<i>a3</i>
8	<i>t0</i>
9	<i>t1</i>
10	<i>t2</i>
11	<i>t3</i>
12	<i>t4</i>
13	<i>t5</i>
14	<i>t6</i>
15	<i>t7</i>
16	<i>s0</i>
17	<i>s1</i>
18	<i>s2</i>
19	<i>s3</i>
20	<i>s4</i>
21	<i>s5</i>
22	<i>s6</i>
23	<i>s7</i>
24	<i>t8</i>

25	<i>t9</i>
26	<i>k0</i>
27	<i>k1</i>
28	<i>gp</i>
29	<i>sp</i>
30	<i>fp</i>
31	<i>ra</i>

The #include command

Source files can contain the `#include filename` command, which has the effect of putting in place of the command row the content of the file *filename*. It is useful if you want to include external routines, and it comes with a loop-detection algorithm that will warn you if you try to do something like “`#include A.s`” in file *B.s* and “`#include B.s`” in file *A.s*.

The instruction set

In this section we will the subset of the MIPS64 instruction set that EduMIPS64 recognizes. We can operate two different taxonomic classification: one based on the functionality of the instructions and one based on the type of the parameters of the instructions.

The first classification divides instruction into three categories: ALU instructions, Load/Store instructions, Flow control instructions. The next three subsections will describe each category and every instruction that belongs to those categories.

The fourth subsection will describe instructions that do not fit in any of the three categories.

ALU Instructions

The Arithmetic Logic Unit (in short, ALU) is a part of the execution unit of a CPU, that has the duty of doing arithmetical and logic operations. So in the ALU instructions group we will find those instructions that do this kind of operations.

ALU Instructions can be divided in two groups: *R-Type* and *I-Type*.

Four of those instructions make use of two special registers: LO and HI. They are internal CPU registers, whose value can be accessed through the *MFLO* and *MFHI* instructions.

Here's the list of R-Type ALU Instructions.

- *AND rd, rs, rt*

Executes a bitwise AND between *rs* and *rt*, and puts the result into *rd*.

- *ADD rd, rs, rt*

Sums the content of 32-bits registers *rs* and *rt*, considering them as signed values, and puts the result into *rd*. If an overflow occurs then trap.

- *ADDU rd, rs, rt*

Sums the content of 32-bits registers *rs* and *rt*, and puts the result into *rd*. No integer overflow occurs under any circumstances.

- *DADD rd, rs, rt*

Sums the content of 64-bits registers *rs* and *rt*, considering them as signed values, and puts the result into *rd*. If an overflow occurs then trap.

- *DADDU rd, rs, rt*

Sums the content of 64-bits registers *rs* and *rt*, and puts the result into *rd*. No integer overflow occurs under any circumstances.

- *DDIV rs, rt*

Executes the division between 64-bits registers *rs* and *rt*, putting the 64-bits quotient in *LO* and the 64-bits remainder in *HI*.

- *DDIVU rs, rt*

Executes the division between 64-bits registers *rs* and *rt*, considering them as unsigned values and putting the 64-bits quotient in *LO* and the 64-bits remainder in *HI*.

- *DIV rs, rt*

Executes the division between 32-bits registers *rs* and *rt*, putting the 32-bits quotient in *LO* and the 32-bits remainder in *HI*.

- *DIVU rs, rt*

Executes the division between 32-bits registers *rs* and *rt*, considering them as unsigned values and putting the 32-bits quotient in *LO* and the 32-bits remainder in *HI*.

- *DMULT rs, rt*

Executes the multiplication between 64-bits registers *rs* and *rt*, putting the low-order 64-bits doubleword of the result into special register *LO* and the high-order 64-bits doubleword of the result into special register *HI*.

- *DMULU rs, rt*

Executes the multiplication between 64-bits registers *rs* and *rt*, considering them as unsigned values and putting the low-order 64-bits doubleword of the result into special register *LO* and the high-order 64-bits doubleword of the result into special register *HI*.

- *DMULTU rs, rt*

Executes the multiplication between 64-bits registers *rs* and *rt*, considering them as unsigned values and putting the low-order 64-bits doubleword of the result into special register *LO* and the high-order 64-bits doubleword of the result into special register *HI*.

- *DSLL rd, rt, sa*

Does a left shift of 64-bits register *rt*, by the amount specified in the immediate (positive) value *sa*, and puts the result into 64-bits register *rd*. Empty bits are padded with zeros.

- *DSLLV rd, rt, rs*

Does a left shift of 64-bits register *rt*, by the amount specified in low-order 6-bits of *rs* threatd as unsigned value, and puts the result into 64-bits register *rd*. Empty bits are padded with zeros.

- *DSRA rd, rt, sa*

Does a right shift of 64-bits register *rt*, by the amount specified in the immediate (positive) value *sa*, and puts the result into 64-bits register *rd*. Empty bits are padded with zeros if the leftmost bit of *rt* is zero, otherwise they are padded with ones.

- *DSRAV rd, rt, rs*

Does a right shift of 64-bits register *rt*, by the amount specified in low-order 6-bits of *rs* threatd as unsigned value, and puts the result into 64-bits register *rd*. Empty bits are padded with zeros if the leftmost bit of *rt* is zero, otherwise they are padded with ones.

- *DSRL rd, rs, sa*

Does a right shift of 64-bits register *rs*, by the amount specified in the immediate (positive) value *sa*, and puts the result into 64-bits register *rd*. Empty bits are padded with zeros.

- *DSRLV rd, rt, rs*

Does a right shift of 64-bits register *rt*, by the amount specified in low-order 6-bits of *rs* threatd as unsigned value, and puts the result into 64-bits register *rd*. Empty bits are padded with zeros.

- *DSUB rd, rs, rt*

Subtracts the value of 64-bits register *rt* to 64-bits register *rs*, considering them as signed values, and puts the result in *rd*. If an overflow occurs then trap.

- *DSUBU rd, rs, rt*

Subtracts the value of 64-bits register *rt* to 64-bits register *rs*, and puts the result in *rd*. No integer overflow occurs under any circumstances.

- *MFLO rd*

Moves the content of the special register LO into *rd*.

- *MFHI rd*

Moves the content of the special register HI into *rd*.

- *MOVN rd, rs, rt*

If *rt* is different from zero, then moves the content of *rs* into *rd*.

- *MOVZ rd, rs, rt*

If *rt* is equal to zero, then moves the content of *rs* into *rd*.

- *MULT rs, rt*

Executes the multiplication between 32-bits registers *rs* and *rt*, putting the low-order 32-bits word of the result into special register LO and the high-order 32-bits word of the result into special register HI.

- *MULTU rs, rt*

Executes the multiplication between 32-bits registers *rs* and *rt*, considering them as unsigned values and putting the low-order 32-bits word of the result into special register LO and the high-order 32-bits word of the result into special register HI.

- *OR rd, rs, rt*

Executes a bitwise OR between *rs* and *rt*, and puts the result into *rd*.

- *SLL rd, rt, sa*

Does a left shift of 32-bits register *rt*, by the amount specified in the immediate (positive) value *sa*, and puts the result into 32-bits register *rd*. Empty bits are padded with zeros.

- *SLLV rd, rt, rs*

Does a left shift of 32-bits register *rt*, by the amount specified in low-order 5-bits of *rs* threatd as unsigned value, and puts the result into 32-bits register *rd*. Empty bits are padded with zeros.

- *SRA rd, rt, sa*

Does a right shift of 32-bits register *rt*, by the amount specified in the immediate (positive) value *sa*, and puts the result into 32-bits register *rd*. Empty bits are padded with zeros if the leftmost bit of *rt* is zero, otherwise they are padded with ones.

- *SRAV rd, rt, rs*

Does a right shift of 32-bits register *rt*, by the amount specified in low-order 5-bits of *rs* threatd as unsigned value, and puts the result into 32-bits register *rd*. Empty bits are padded with zeros if the leftmost bit of *rt* is zero, otherwise they are padded with ones.

- *SRL rd, rs, sa*

Does a right shift of 32-bits register *rs*, by the amount specified in the immediate (positive) value *sa*, and puts the result into 32-bits register *rd*. Empty bits are padded with zeros.

- *SRLV rd, rt, rs*

Does a right shift of 32-bits register *rt*, by the amount specified in low-order 5-bits of *rs* threatd as unsigned value, and puts the result into 32-bits register *rd*. Empty bits are padded with zeros.

- *SUB rd, rs, rt*

Subtracts the value of 32-bits register *rt* to 32-bits register *rs*, considering them as signed values, and puts the result in *rd*. If an overflow occurs then trap.

- *SUBU rd, rs, rt*

Subtracts the value of 32-bits register *rt* to 32-bits register *rs*, and puts the result in *rd*. No integer overflow occurs under any circumstances.

- *SLT rd, rs, rt*

Sets the value of *rd* to 1 if the value of *rs* is less than the value of *rt*, otherwise sets it to 0. This instruction performs a signed comparison.

- *SLTU rd, rs, rt*

Sets the value of *rd* to 1 if the value of *rs* is less than the value of *rt*, otherwise sets it to 0. This instruction performs an unsigned comparison.

- *XOR rd, rs, rt*

Executes a bitwise exclusive OR (XOR) between *rs* and *rt*, and puts the result into *rd*.

Here's the list of I-Type ALU Instructions.

- *ADDI rt, rs, immediate*

Executes the sum between 32-bits register *rs* and the immediate value, putting the result in *rt*. This instruction considers *rs* and the immediate value as signed values. If an overflow occurs then trap.

- *ADDIU rt, rs, immediate*

Executes the sum between 32-bits register *rs* and the immediate value, putting the result in *rt*. No integer overflow occurs under any circumstances.

- *ANDI rt, rs, immediate*

Executes the bitwise AND between *rs* and the immediate value, putting the result in *rt*.

- *DADDI rt, rs, immediate*

Executes the sum between 64-bits register *rs* and the immediate value, putting the result in *rt*. This instruction considers *rs* and the immediate value as signed values. If an overflow occurs then trap.

- *DADDIU rt, rs, immediate*

Executes the sum between 64-bits register *rs* and the immediate value, putting the result in *rt*. No integer overflow occurs under any circumstances.

- *DADDUI rt, rs, immediate*

Executes the sum between 64-bits register *rs* and the immediate value, putting the result in *rt*. No integer overflow occurs under any circumstances.

- *LUI rt, immediate*

Loads the constant defined in the immediate value in the upper half (16 bit) of the lower 32 bits of *rt*, sign-extending the upper 32 bits of the register.

- *ORI rt, rs, immediate*

Executes the bitwise OR between *rs* and the immediate value, putting the result in *rt*.

- *SLTI rt, rs, immediate*

Sets the value of *rt* to 1 if the value of *rs* is less than the value of the immediate, otherwise sets it to 0. This instruction performs a signed comparison.

- *SLTUI rt, rs, immediate*

Sets the value of *rt* to 1 if the value of *rs* is less than the value of the immediate, otherwise sets it to 0. This instruction performs an unsigned comparison.

- *XORI rt, rs, immediate*

Executes a bitwise exclusive OR (XOR) between *rs* and the immediate value, and puts the result into *rt*.

Load/Store instructions

This category contains all the instructions that operate transfers between registers and the memory. All of these instructions are in the form:

```
[label:] instruction rt, offset(base)
```

Where *rt* is the source or destination register, depending if we are using a store or a load instruction; *offset* is a label or an immediate value and *base* is a register. The address is obtained by adding to the value of the register *base* the immediate value *offset*.

The address specified must be aligned according to the data type that is treated. Load instructions ending with “U” treat the content of the register *rt* as an unsigned value.

List of load instructions:

- *LB rt, offset(base)*

Loads the content of the memory cell at address specified by *offset* and *base* in register *rt*, treating it as a signed byte.

- *LBU rt, offset(base)*

Loads the content of the memory cell at address specified by *offset* and *base* in register *rt*, treating it as an unsigned byte.

- *LD rt, offset(base)*

Loads the content of the memory cell at address specified by *offset* and *base* in register *rt*, treating it as a double word.

- *LH rt, offset(base)*

Loads the content of the memory cell at address specified by *offset* and *base* in register *rt*, treating it as a signed half word.

- *LHU rt, offset(base)*

Loads the content of the memory cell at address specified by *offset* and *base* in register *rt*, treating it as an unsigned half word.

- *LW rt, offset(base)*

Loads the content of the memory cell at address specified by *offset* and *base* in register *rt*, treating it as a signed word.

- *LWU rt, offset(base)*

Loads the content of the memory cell at address specified by *offset* and *base* in register *rt*, treating it as a signed word.

List of store instructions:

- *SB rt, offset(base)*

Stores the content of register *rt* in the memory cell specified by *offset* and *base*, treating it as a byte.

- *SD rt, offset(base)*

Stores the content of register *rt* in the memory cell specified by *offset* and *base*, treating it as a double word.

- *SH rt, offset(base)*

Stores the content of register *rt* in the memory cell specified by *offset* and *base*, treating it as a half word.

- *SW rt, offset(base)*

Stores the content of register *rt* in the memory cell specified by *offset* and *base*, treating it as a word.

Flow control instructions

Flow control instructions are used to alter the order of instructions that are fetched by the CPU. We can make a distinction between these instructions: R-Type, I-Type and J-Type.

Those instructions effectively executes the jump in the ID stage, so often an useless fetch is executed. In this case, two instructions are removed from the pipeline, and the branch taken stalls counter is incremented by two units.

List of R-Type flow control instructions:

- *JALR rs*

Puts the content of *rs* into the program counter, and puts into R31 the address of the instruction that follows the JALR instruction, the return value.

- *JR rs*

Puts the content of *rs* into the program counter.

List of I-Type flow control instructions:

- *B offset*

Unconditionally jumps to *offset*

- *BEQ rs, rt, offset*

Jumps to *offset* if *rs* is equal to *rt*.

- *BEQZ rs, offset*

Jumps to *offset* if *rs* is equal to zero.

- *BGEZ rs, offset*

If *rs* is greater than zero, does a PC-relative jump to *offset*.

- *BNE rs, rt, offset*

Jumps to *offset* if *rs* is not equal to *rt*.

- *BNEZ rs, offset*

Jumps to *offset* if *rs* is not equal to zero.

List of J-Type flow control instructions:

- *J target*

Puts the immediate value *target* into the program counter.

- *JAL target*

Puts the immediate value *target* into the program counter, and puts into R31 the address of the instruction that follows the JAL instruction, the return value.

The SYSCALL instruction

The SYSCALL instruction offers to the programmer an operating-system-like interface, making available six different system calls.

System calls expect that the address of their parameters is stored in register R14 (\$t6), and will put their return value in register R1 (\$at).

System calls follow as much as possible the POSIX convention.

SYSCALL 0 - exit()

SYSCALL 0 does not expect any parameter, nor it returns anything. It simply stops the simulator.

Note that if the simulator does not find SYSCALL 0 in the source code, or any of its equivalents (HALT - TRAP 0), it will be added automatically at the end of the source.

SYSCALL 1 - open()

The SYSCALL 1 expects two parameters: a zero-terminated string that indicates the pathname of the file that must be opened, and a double word containing an integer that indicates the flags that must be used to specify how to open the file.

This integer must be built summing the flags that you want to use, choosing them from the following list:

- *O_RDONLY (0x01)* Opens the file in read only mode;
- *O_WRONLY (0x02)* Opens the file in write only mode;
- *O_RDWR (0x03)* Opens the file in read/write mode;

- *O_CREAT (0x04)* Creates the file if it does not exist;
- *O_APPEND (0x08)* In write mode, appends written text at the end of the file;
- *O_TRUNC (0x08)* In write mode, deletes the content of the file as soon as it is opened.

It is mandatory to specify one of the first three modes. The fourth and the fifth modes are exclusive, you can not specify *O_APPEND* if you specify *O_TRUNC* (and vice versa).

You can specify a combination of modes by simply adding the integer values of those flags. For instance, if you want to open a file in write only mode and append the written text to the end of file, you should specify the mode $2 + 8 = 10$.

The return value of the system call is the new file descriptor associated with the file, that can be further used with the other system calls. If there is an error, the return value will be -1.

SYSCALL 2 - close()

SYSCALL 2 expects only one parameter, the file descriptor of the file that is closed.

If the operation ends successfully, SYSCALL 2 will return 0, otherwise it will return -1. Possible causes of failure are the attempt to close a non-existent file descriptor or the attempt to close file descriptors 0, 1 or 2, that are associated respectively to standard input, standard output and standard error.

SYSCALL 3 - read()

SYSCALL 3 expects three parameters: the file descriptor to read from, the address where the read data must be put into, the number of bytes to read.

If the first parameter is 0, the simulator will prompt the user for an input, via an input dialog. If the length of the input is greater than the number of bytes that have to be read, the simulator will show again the message dialog.

It returns the number of bytes that have effectively been read, or -1 if the read operation fails. Possible causes of failure are the attempt to read from a non-existent file descriptor, the attempt to read from file descriptors 1 (standard output) or 2 (standard error) or the attempt to read from a write-only file descriptor.

SYSCALL 4 - write()

SYSCALL 4 expects three parameters: the file descriptor to write to, the address where the data must be read from, the number of bytes to write.

If the first parameter is two or three, the simulator will pop the input/output frame, and write there the read data.

It returns the number of bytes that have been written, or -1 if the write operation fails. Possible causes of failure are the attempt to write to a non-existent file descriptor, the attempt to write to file descriptor 0 (standard input) or the attempt to write to a read-only file descriptor.

SYSCALL 5 - printf()

SYSCALL 5 expects a variable number of parameters, the first being the address of the so-called "format string". In the format string can be included some placeholders, described in the following list:

- *%s* indicates a string parameter;
- *%i* indicates an integer parameter;
- *%d* behaves like *%i*;
- *%%* literal %

For each *%s*, *%d* or *%i* placeholder, SYSCALL 5 expects a parameter, starting from the address of the previous one.

When the SYSCALL finds a placeholder for an integer parameter, it expects that the corresponding parameter is an integer value, when if it finds a placeholder for a string parameter, it expects as a parameter the address of the string.

The result is printed in the input/output frame, and the number of bytes written is put into R1.

If there's an error, -1 is written to R1.

Other instructions

In this section there are instructions that do not fit in the previous categories.

BREAK

The BREAK instruction throws an exception that has the effect to stop the execution if the simulator is running. It can be used for debugging purposes.

NOP

The NOP instruction does not do anything, and it's used to create gaps in the source code.

TRAP

The TRAP instruction is a deprecated alias for the SYSCALL instruction.

HALT

The HALT instruction is a deprecated alias for the SYSCALL 0 instruction, that halts the simulator.

Floating Point Unit

This chapter ¹ describes the Floating Point Unit (FPU) emulated in EduMIPS64.

In the first paragraph we introduce the double format, the special floating point values defined in the IEEE 754 standard and the exceptions that floating point computations can raise.

In the second paragraph we explain how EduMIPS64 allows users to enable or disable the IEEE floating point traps.

In the third paragraph we describe how double precision numbers and special values can be specified in the source programs.

In the fourth paragraph, we introduce the FCSR register, used by the FPU to represent its state. It contains information about rounding, the boolean results of comparison operations and the policies for handling IEEE floating point exceptions.

In the fifth and last paragraph, we present all the MIPS64 floating point instructions that have been implemented in EduMIPS64.

Before starting the discussion about the FPU, we define the domain of floating point double precision numbers as $[-1.79E308, -4.94E-324] \cup \{0\} \cup [4.94E-324, 1.79E308]$.

Special values

Floating point arithmetics allows the programmer to choose whether to stop the computation or not, if invalid operations are carried on. In this scenario, operations like the division between zeroes or square roots of negative numbers must produce a result that, not being a number (NaN) is treated as something different.

NaN or Invalid Operation

The IEEE Standard for Floating-Point Arithmetic (IEEE 754) defined that invalid arithmetic operations can either signal the error while the program is running (using a trap for the IEEE exception **Invalid Operation**) or return as a result the special value QNaN (Quit Not a Number). Another NaN value, that inconditionally raises the same trap once it is detected as being one of the operands, is SNaN (Signalling Not a Number). This value is seldom used in applications, and historically it has been used to initialize variables.

¹ This chapter is part of the Bachelor's degree thesis by Massimo Trubia: "Progetto e implementazione di un modello di Floating Point Unit per un simulatore di CPU MIPS64".

Zeroes or Underflows

Another special value defined by the standard is zero. Since the double format does not include the zero in its domain, it is considered a special value. There is a positive zero and a negative zero: the former is used when a representation of a negative number $\in]-4.94\text{E-}324, 0[$ is attempted, and a result is required (as opposed to allowing an **Underflow** trap), while the latter is used when the number that should be represented is $\in [0, 4.94\text{E-}324[$, and the Underflow trap is disabled.

Infinities or Overflows

When a program attempts to represent a value with an extremely large absolute value ($\in]-\infty, -1.79\text{E}308[\cup]1.79\text{E}308, +\infty[$), that is outside the domain of double values, the CPU returns either $-\infty$ or $+\infty$. The alternative is to trigger a trap for the exceptional **Overflow** condition.

Infinities can also be returned in case of a division by zero; in that case the sign of the infinite is given by the product of the sign of the zero and the sign of the dividend. The **Divide by zero** trap can be alternatively raised.

Exception configuration

EduMIPS64 allows the user to enable or disable the traps for 4 of the 5 IEEE exceptions, through the *FPU Exceptions* tab in the *Configure* → *Settings* window. If any of them is disabled, the respective special value will be returned (as described in Special values).

The .double directive

The `.double` directive must be used in the `.data` section of source files, and allows to allocate a memory cell for a *double* value.

The directive can be used in 2 ways:

```
variable-name: .double double_number
variable-name: .double keyword
```

where `double_number` can be represented either in extended notation (`1.0,0.003`), or in scientific notation (`3.7E-12,0.5E32`). `keyword` can be `POSITIVEINFINITY`, `NEGATIVEINFINITY`, `POSITIVEZERO`, `NEGATIVEZERO`, `SNAN` e `QNAN`, thus allowing to directly insert in memory the special values.

The FCSR register

The FCSR (Floating point Control Status Register) is the register that controls several functional aspects of the FPU. It is 32 bits long and it is represented in the statistics window.

The **FCC** field is 8 bits wide, from 0 to 7. The conditional instructions (`C.EQ.D`, `C.LT.D`) use it to save the boolean result of comparisons between two registers.

The Cause, Enables and Flag fields handle the dynamics of IEEE exceptions described in Special values. Each of them is composed of 5 bits, V (Invalid Operation), Z (Divide by Zero), O (Overflow), U (Underflow) and I (Inexact); the latter is not yet used.

The **Clause** field bits are set if the corresponding IEEE exceptions occur during the execution of a program.

The **Enable** field bits are set through the configuration window and show the IEEE exceptions for which traps are enabled.

The **Flag** field shows the exceptions that have happened but, since the trap is not enabled for that particular exception, have returned special values (the ones described in Special values).

The **RM** field describes the rounding method currently in use to convert floating point numbers to integers (see the description of the `CVT.L.D` instruction).

Instruction set

This section describes the MIPS64 FPU instruction implemented by EduMIPS64; they are listed in alphabetic order. The operations performed by the instruction are described using a notation according to which the *i*-th memory cell is represented as `memory[i]`, and the FCC fields of the FCSR register are `FCSR_FCC[cc]`, `cc` $\in [0,7]$.

In some instructions, to avoid ambiguity, the registers are represented as `GPR[i]` and `FPR[i]`, $i \in [0,31]$, but in most cases we just use the `rx` or `fx` notation, with $x \in \{d,s,t\}$. The three letters are used to indicate the purpose of each register (destination, source, third). Lastly, the values returned by conversion operations are represented with the following notation: `convert_conversiontype(register[,rounding_type])`, where the `rounding_type` parameter is optional.

Some examples for the FPU instructions are available at <http://www.edumips.org/attachment/wiki/Upload/FPUMaxSamples.rar>.

- **ADD.D fd, fs, ft**

Description: `fd = fs + ft`

Exceptions: Overflow and underflow traps are generated if the result cannot be represented according to IEEE 754. Invalid operation is raised if `fs` or `ft` contain QNaN or SNaN, or if an invalid operation ($+\infty - \infty$) is executed.

- **BC1F cc, offset**

Description: if `FCSR_FCC[cc] == 0` then branch

If `FCSR_FCC[cc]` is false, do a PC-relative branch.

Example:

```
C.EQ.D 7,f1,f2
BC1F 7,label
```

In this example, `C.EQ.D` checks if `f1` and `f2` are equal, writing the results of the comparison in the 7th bit of the FCC field of the FCSR register. After that, `BC1F` jumps to `label` if the result of the comparison is 0 (false).

- **BC1T cc, offset**

Description: if `FCSR_FCC[cc] == 1` then branch

If `FCSR_FCC[cc]` is true, do a PC-relative branch.

Example:

```
C.EQ.D 7,f1,f2
BC1T 7,label
```

In this example, `C.EQ.D` checks if `f1` and `f2` are equal, writing the results of the comparison in the 7th bit of the FCC field of the FCSR register. After that, `BC1T` jumps to `label` if the result of the comparison is 1 (true).

- **C.EQ.D cc, fs, ft**

Description: `FCSR_FCC[cc] = (fs==ft)`

Checks if `fs` is equal to `ft`, and saves the result of the comparison in `FCSR_FCC[cc]`. See examples for `BC1T`, `BC1F`.

Exceptions: Invalid Operation can be thrown if `fs` or `ft` contain QNaN (trap is triggered if it is enabled) or SNaN (trap is always triggered).

- **C.LT.D cc, fs, ft**

Description: `FCSR_FCC[cc] = (fs<ft)`

Checks if `fs` is smaller than `ft`, and saves the result of the comparison in `FCSR_FCC[cc]`.

Example:

```
C.LT.D 2,f1,f2
BC1T 2,target
```

In this example, `C.LT.D` checks if `f1` is smaller than `f2`, and saves the result of the comparison in the second bit of the FCC field of the FCSR register. After that, `BC1T` jumps to `target` if that bit is set to 1.

Exceptions: Invalid Operation can be thrown if `fs` or `ft` contain QNaN (trap is triggered if it is enabled) or SNaN (trap is always triggered).

- **CVT.D.L fd, fs**

Description: `fd = convert_longToDouble(fs)`

Converts a long to a double.

Example:

```
DMTC1 r6,f5
CVT.D.L f5,f5
```

In this example, `DMTC1` copies the value of GPR `r6` to FPR `f5`; after that `CVT.D.L` converts the value stored in `f5` from long to double. If for instance `r6` contains the value 52, after the execution of `DMTC1` the binary representation of 52 gets copied to `f5`. After the execution of `CVT.D.L`, `f5` contains the IEEE 754 representation of 52.0.

Exceptions: Invalid Operation is thrown if `fs` contains QNaN, SNaN or an infinite.

- **CVT.D.W `fd, fs`**

Description: `fd = convert_IntToDouble(fs)`

Converts an int to a double.

Example:

```
MTC1 r6,f5
CVT.D.W f5,f5
```

In this example, `MTC1` copies the lower 32 bit of the GPR `r6` into the FPR `f5`. Then, `CVT.D.W`, reads `f5` as an int, and converts it to double.

If we had `r6=0xAAAAAAAABBBBBBBB`, after the execution of `MTC1` we get `f5=0xFFFFFFFFBBBBBBBB`; its upper 32 bits (XX..X) are now UNDEFINED (haven't been overwritten). `CVT.D.W` interprets `f5` as an int (`f5=-1145324613`), and converts it to double (`f5=0xC1D1111111400000 = -1.145324613E9`).

Exceptions: Invalid Operation is thrown if `fs` contains QNaN, SNaN or an infinite.

- **CVT.L.D `fd, fs`**

Description: `fd = convert_doubleToLong(fs, CurrentRoundingMode)`

Converts a double to a long, rounding it before the conversion.

Example:

```
CVT.L.D f5,f5
DMFC1 r6,f5
```

`CVT.L.D` the double value in `f5` to a long; then `DMFC1` copies `f5` to `r6`; the result of this operation depends on the current rounding modality, that can be set in the *FPU Rounding* tab of the *Configure* → *Settings* window.

Rounding examples

Tipo	RM field	f5 register	r6 register
To nearest	0	6.4	6
To nearest	0	6.8	7
To nearest	0	6.5	6 (to even)
To nearest	0	7.5	8 (to even)
Towards 0	1	7.1	7
Towards 0	1	-2.3	-2
Towards ∞	2	4.2	5
Towards ∞	2	-3.9	-3
Towards $-\infty$	3	4.2	4

Towards $-\infty$	3	-3.9	-4
-------------------	---	------	----

- *CVT.W.D fd, fs*

Description: `fd = convert_DoubleToInt(fs, CurrentRoundingMode)`

Converts a double to an int, using the current rounding modality.

Exceptions: Invalid Operation is thrown if fs contains an infinite value, any NaN or the results is outside the signed int domain $[-2^{63}, 2^{63}-1]$

- *DIV.D fd, fs, ft*

Description: `fd = fs \div ft`

Exceptions: Overflow or Underflow are raised if the results cannot be represented using the IEEE 754 standard. Invalid Operation is raised if fs or ft contain QNaN or SNaN, or if an invalid operation is executed ($0 \div 0, \infty \div \infty$). Divide by zero is raised if a division by zero is attempted with a dividend that is not QNaN or SNaN.

- *DMFC1 rt, fs*

Description: `rt = fs`

Executes a bit per bit copy of the FPR fs into the GPR rt.

- *DMTC1 rt, fs*

Description: `fs = rt`

Executes a bit per bit copy of the GPR rt into the FPR fs.

- *L.D ft, offset(base)*

Description: `ft = memory[GPR[base] + offset]`

Loads from memory a doubleword and stores it in ft.

Note

L.D is not present in the MIPS64 ISA, it is an alias for *LDC1* that is present in EduMIPS64 for compatibility with WinMIPS64.

- *LDC1 ft, offset(base)*

Description: `memory[GPR[base] + offset]`

Loads from memory a doubleword and stores it in ft.

- *LWC1 ft, offset(base)*

Description: `ft = memory[GPR[base] + offset]`

Loads from memory a word and stores it in ft.

- *MFC1 rt, fs*

Description: `rt = readInt(fs)`

Reads the fs FPR as an int and writes its value to the rt GPR as long. *Example:*

```
MFC1 r6, f5
SD r6, mem(R0)
```

Let `f5=0xAAAAAAAABBBBBBBB`; *MFC1* reads `f5` as an int (lower 32 bits), interpreting `BBBBBBBB` as `-1145324613`, and writes the value to `f6` (64 bits). After the execution of *MFC1*, `r6=0xFFFFFFFFBBBBBBBB` = `-1145324613`. So the *SD* instruction will write to memory a doubleword with this value, since the sign in `r6` was extended.

- *MOV.F.D fd, fs, cc*

Description: if `FCSR_FCC[cc] == 0` then `fd=fs`

If `FCSR_FCC[cc]` is false, the copies `fs` to `fd`.

- *MOVT.D fd, fs, cc*

Description: if `FCSR_FCC[cc] == 1` then `fd=fs`

If `FCSR_FCC[cc]` is true, the copies `fs` to `fd`.

- *MOV.D fd, fs*

Description: `fd = fs`

Copies `fs` to `fd`.

- *MOVN.D fd, fs, rt*

Description: if `rt != 0` then `fd=fs`

If `rt` is not zero, copies `fs` to `fd`.

- *MOVZ.D fd, fs, rt*

Description: if `rt == 0` then `fd=fs`

If `rt` is equal to zero, copies `fs` to `fd`.

- *MTC1 rt, fs*

Description: `fs = rt0..31`

Copies the lower 32 bit of `rt` to `fs`.

Example:

```
MTC1 r6, f5
```

Let `r5=0xAAAAAAAABBBBBBBB`; `MTC1` reads the lower 32 bits of `r5` copying them to the 32 lower bits of `f5`. The higher 32 bits of `f5` are not overwritten.

- *MUL.D fd, fs, ft*

Description: `fd = fs × ft`

Exceptions: Overflow or Underflow are raised if the results cannot be represented using the IEEE 754 standard. Invalid Operation is raised if `fs` or `ft` contain QNaN or SNaN, or if an invalid operation is executed (multiply by ∞ OR BY QNaN).

- *S.D ft, offset(base)*

Description: `memory[base+offset] = ft`

Copies `ft` to memory.

Note

`S.D` is not present in the MIPS64 ISA, it is an alias for `SDC1` that is present in EduMIPS64 for compatibility with WinMIPS64.

- *SDC1 ft, offset(base)*

Description: `memory[base+offset] = ft`

Copies `ft` to memory.

- *SUB.D fd, fs, ft*

Description: `fd = fs - ft`

Exceptions: Overflow and underflow traps are generated if the result cannot be represented according to IEEE 753. Invalid operation is raised if `fs` or `ft` contain QNaN or SNaN, or if an invalid operation ($+\infty - \infty$) is executed.

- *SWC1 ft, offset(base)*

Description: `memory[base+offset] = ft`

Copies the lower 32 bits of `ft` to memory.

The user interface

The GUI of EduMIPS64 is inspired to WinMIPS64 user interface. In fact, the main window is identical, except for some menus.

The EduMIPS64 main window is composed by a menu bar and six frames, showing different aspects of the simulation. There's also a status bar, that has the double purpose to show the content of memory cells and registers when you click them and to notify the user that the simulator is running when the simulation has been started but verbose mode is not selected.

The status bar also shows the CPU status. It can show one of the following four states:

- *READY* The CPU hasn't executed any instructions (no program is loaded).
- *RUNNING* The CPU is executing a series of instructions.
- *STOPPING* The CPU has found a termination instruction, and is executing the instructions that are already in the pipeline before terminating the execution.
- *HALTED* The CPU is stopped: a program just finished running.

Note that the CPU status is different from the simulator status. The simulator may execute a number of CPU cycles and then stop executing, allowing the user to inspect memory and registers: in this state, between CPU cycles, the CPU stays in *RUNNING* or *STOPPING* state. Once the CPU reaches the *HALTED* state, the user cannot run any CPU cycle without loading a program again (the same program, or a different one).

There are more details in the following sections.

The menu bar

The menu bar contains six menus:

File

The File menu contains menu items about opening files, resetting or shutting down the simulator, writing trace files.

- *Open...* Opens a dialog that allows the user to choose a source file to open.
- *Open recent* Shows the list of the recent files opened by the simulator, from which the user can choose the file to open
- *Reset* Resets the simulator, keeping open the file that was loaded but resetting the execution.
- *Write Dinero Tracefile...* Writes the memory access data to a file, in xdin format.
- *Exit* Closes the simulator.

The *Write Dinero Tracefile...* menu item is only available when a whole source file has been executed and the end has been already reached.

Execute

The Execute menu contains menu items regarding the execution flow of the simulation.

- *Single Cycle* Executes a single simulation step
- *Run* Starts the execution, stopping when the simulator reaches a *SYSCALL 0* (or equivalent) or a *BREAK* instruction, or when the user clicks the Stop menu item (or presses F9).
- *Multi Cycle* Executes some simulation steps. The number of steps executed can be configured through the Setting dialog.
- *Stop* Stops the execution when the simulator is in "Run" or "Multi cycle" mode, as described previously.

This menu is only available when a source file is loaded and the end of the simulation is not reached. The *Stop* menu item is available only in "Run" or "Multi Cycle" mode.

Note that the simulator slows down when updating the UI. If you want to execute long (thousands of cycles) programs quickly, disable the "Sync graphics with CPU in multi-step execution" option.

Configure

The Configure menu provides facilities for customizing EduMIPS64 appearance and behavior.

- *Settings...* Opens the Settings dialog, described in the next sections of this chapter;
- *Change Language* Allows the user to change the language used by the user interface. Currently only English and Italian are supported. This change affects every aspect of the GUI, from the title of the frames to the online manual and warning/error messages.

The *Settings...* menu item is not available when the simulator is in “Run” or “Multi Cycle” mode, because of potential race conditions.

Tools

This menu contains only an item, used to invoke the Dinero Frontend dialog.

- *Dinero Frontend...* Opens the Dinero Frontend dialog.

This menu is not available until you have not executed a program and the execution has reached its end.

Window

This menu contains items related to operations with frames.

- *Tile* Sorts the visible windows so that no more than three frames are put in a row. It tries to maximize the space occupied by every frame.

The other menu items simply toggle the status of each frame, making them visible or minimizing them.

Help

This menu contains help-related menu items.

- *Manual...* Shows the Help dialog.
- *About us...* Shows a cute dialog that contains the names of the project contributors, along with their roles.

Frames

The GUI is composed by seven frames, six of which are visible by default, and one (the I/O frame) is hidden.

Cycles

The Cycles frame shows the evolution of the execution flow during time, showing for each time slot which instructions are in the pipeline, and in which stage of the pipeline they are located.

Registers

The Registers frame shows the content of each register. By left-clicking on them you can see in the status bar their decimal (signed) value, while double-clicking on them will pop up a dialog that allows the user to change the value of the register.

Statistics

The Statistics frame shows some statistics about the program execution.

Note that during the last execution cycle the cycles counter is not incremented, because the last execution cycle is not a full CPU cycle but rather a pseudo-cycle whose only duties are to remove the last instruction from the pipeline and increment the counter of executed instructions.

Pipeline

The Pipeline frame shows the actual status of the pipeline, showing which instruction is in which pipeline stage. Different colors highlight different pipeline stages.

Memory

The Memory frame shows memory cells content, along with labels and comments taken from the source code. Memory cells content, like registers, can be modified double-clicking on them, and clicking on them will show their decimal value in the status bar. The first column shows the hexadecimal address of the memory cell, and the second column shows the value of the cell. Other columns show additional info from the source code.

Code

The Code window shows the instructions loaded in memory. The first column shows the address of the instruction, while the second column shows the hexadecimal representation of the instructions. Other columns show additional info taken from the source code.

Input/Output

The Input/Output window provides an interface for the user to see the output that the program creates through the SYSCALLs 4 and 5. Actually it is not used for input, as there's a dialog that pops up when a SYSCALL 3 tries to read from standard input, but future versions will include an input text box.

Dialogs

Dialogs are used by EduMIPS64 to interact with the user in many ways. Here's a summary of the most important dialogs:

Settings

In the Settings dialog various aspects of the simulator can be configured. Clicking on the "OK" button will cause the options to be saved, while clicking on "Cancel" (or simply closing the window) will cause the changes to be ignored. Don't forget to click "OK" if you want to save your changes.

The Main Settings tab allow to configure forwarding and the number of steps in the Multi Cycle mode.

The Behavior tab allow to enable or disable warnings during the parsing phase, the "Sync graphics with CPU in multi-step execution" option, when checked, will synchronize the frames' graphical status with the internal status of the simulator. This means that the simulation will be slower, but you'll have an explicit graphical feedback of what is happening during the simulation. If this option is checked, the "Interval between cycles" option will influence how many milliseconds the simulator will wait before starting a new cycle. Those options are effective only when the simulation is run using the "Run" or the "Multi Cycle" options from the Execute menu.

The last two options set the behavior of the simulator when a synchronous exception is raised. If the "Mask synchronous exceptions" option is checked, the simulator will ignore any Division by zero or Integer overflow exception. If the "Terminate on synchronous exception" option is checked, the simulation will be halted if a synchronous exception is raised. Please note that if synchronous exceptions are masked, nothing will happen, even if the termination option is checked. If exceptions are not masked and the termination option is not checked, a dialog will pop out, but the simulation will go on as soon as the dialog is closed. If exceptions are not masked and the termination option is checked, the dialog will pop out, and the simulation will be stopped as soon as the dialog is closed.

The last tab allows to change the appearance of the user interface. There are options to change the colors associated to the different pipeline stages, an option to choose whether memory cells are shown as long or double values and an option to set the UI font size.

Note that the UI scaling with font size is far from perfect, but it should be enough to make the simulator usable with high-resolution displays (e.g., 4k).

Dinero Frontend

The Dinero Frontend dialog allows to feed a DineroIV process with the trace file internally generated by the execution of the program. In the first text box there is the path of the DineroIV executable, and in the second one there must be the parameters of DineroIV.

The lower section contains the output of the DineroIV process, from which you can take the data that you need.

Help

The Help dialog brings up the on-line manual, which is an HTML copy of this document.

Command line options

Three command line options are available. They are described in the following list, with the long name enclosed in round brackets. Long and short names can be used in the same way.

- `-v` (`-version`) prints the simulator version and exits.
- `-h` (`-help`) prints a help message with a brief summary of command line options, then exits.
- `-f` (`-file`) *filename* opens *filename* in the simulator
- `-r` (`-reset`) resets the stored configuration to the default values
- `-d` (`-debug`) enters Debug mode

The `-debug` flag has the effect to activate Debug mode. In this mode, a new frame is available, the Debug frame, and it shows the log of internal activities of EduMIPS64. It is not useful for the end user, it is meant to be used by EduMIPS64 developers.

Running EduMIPS64

The EduMIPS64 *.jar* file can act both as a stand-alone executable *.jar* file and as an applet, so it can be executed in both ways. Both methods need the Java Runtime Environment, version 5 or later.

To run it as a stand-alone application, the *java* executable must be issued in this way: *java -jar edumips64-version.jar*, where the *version* string must be replaced with the actual version of the simulator. On some systems, you may be able to execute it by just clicking on the *.jar* file.

To embed it in an HTML, the `<applet>` tag must be used. The EduMIPS64 web site contains a page that already contains the applet, so that everyone can execute it without the hassle of using the command line.

Code Examples

In this chapter you'll find some sample listings that will be useful in order to understand how EduMIPS64 works.

SYSCALL

It's important to understand that examples for SYSCALL 1-4 refer to the *print.s* file, that is the example for SYSCALL 5. If you want to run the examples, you should copy the content of that example in a file named *print.s* and include it in your code.

Some examples use an already existing file descriptor, even if it doesn't truly exist. If you want to run those examples, use the SYSCALL 1 example to open a file.

SYSCALL 0

When SYSCALL 0 is called, it stops the execution of the program. Example:

```
.code
daddi    r1, r0, 0      ; saves 0 in R1
syscall 0                ; exits
```

SYSCALL 1

Example program that opens a file:

```
error_op:      .data
               .asciiz    "Error opening the file"
ok_message:    .asciiz    "All right"
params_sys1:   .asciiz    "filename.txt"
               .word64     0xF
```



```

open:          .text
              daddi      r14, r0, params_sys1
              syscall    1
              daddi      $s0, r0, -1
              dadd       $s2, r0, r1
              daddi      $a0, r0, ok_message
              bne        r1, $s0, end
              daddi      $a0, r0, error_op

end:           jal        print_string
              syscall    0

              #include    print.s

```

In the first two rows we write to memory the strings containing the error message and the success message that we will pass to print_string function, and we give them two labels. The print_string function is included in the print.s file.

Next, we write to memory the data required from SYSCALL 1 (row 4, 5), the path of the file to be opened (that must exist if we work in read or read/write mode) and, in the next memory cell, an integer that defines the opening mode.

In this example, the file was opened using the following modes: `O_RDWR` | `O_CREAT` | `O_APPEND`. The number 15 (0xF in base 16) comes from the sum of the values of these three modes (3 + 4 + 8).

We give a label to this data so that we can use it later.

In the .text section, we save the address of params_sys1 (that for the compiler is a number) in register r14; next we can call SYSCALL 1 and save the content of r1 in \$s2, so that we can use it in the rest of the program (for instance, with other SYSCALL).

Then the print_string function is called, passing error_op as an argument if r1 is equal to -1 (rows 13-14) or else passing ok_message as an argument if everything went smoothly (rows 12 and 16).

SYSCALL 2

Example program that closes a file:

```

params_sys2:   .data
              .space 8
error_cl:      .asciiz    "Error closing the file"
ok_message:    .asciiz    "All right"

close:         .text
              daddi      r14, r0, params_sys2
              sw          $s2, params_sys2(r0)
              syscall    2
              daddi      $s0, r0, -1
              daddi      $a0, r0, ok_message
              bne        r1, $s0, end
              daddi      $a0, r0, error_cl

end:           jal        print_string
              syscall    0

              #include    print.s

```

First we save some memory for the only argument of SYSCALL 2, the file descriptor of the file that must be closed (row 2), and we give it a label so that we can access it later.

Next we put in memory the strings containing the error message and the success message, that will be passed to the print_string function (rows 3, 4).

In the .text section, we save the address of params_sys2 in r14; then we can call SYSCALL 2.

Now we call the `print_string` function using `error_cl` as a parameter if `r1` yields `-1` (row 13), or we call it using `ok_message` as a parameter if all went smoothly (row 11).

Note: This listing needs that registry `$s2` contains the file descriptor of the file to use.

SYSCALL 3

Example program that reads 16 bytes from a file and saves them to memory:

```
.data
params_sys3:  .space      8
ind_value:    .space      8
               .word64    16
error_3:      .asciiz     "Error while reading from file"
ok_message:   .asciiz     "All right"

value:        .space      30

.text
read:         daddi        r14, r0, params_sys3
               sw          $s2, params_sys3(r0)
               daddi        $s1, r0, value
               sw          $s1, ind_value(r0)
               syscall      3
               daddi        $s0, r0, -1
               daddi        $a0, r0, ok_message
               bne          r1, $s0, end
               daddi        $a0, r0, error_3

end:          jal          print_string
               syscall      0

#include      print.s
```

The first 4 rows of the `.data` section contain the arguments of SYSCALL 3, the file descriptor of the from which we must read, the memory address where the SYSCALL must save the read data, the number of bytes to read. We give labels to those parameters that must be accessed later. Next we put, as usual, the strings containing the error message and the success message.

In the `.text` section, we save the `params_sys3` address to register `r14`, we save in the memory cells for the SYSCALL parameters the file descriptor (that we suppose to have in `$s2`) and the address that we want to use to save the read bytes.

Next we can call SYSCALL 3, and then we call the `print_string` function passing as argument `error_3` or `ok_message`, according to the success of the operation.

SYSCALL 4

Example program that writes to a file a string:

```
.data
params_sys4:  .space      8
ind_value:    .space      8
               .word64    16
error_4:      .asciiz     "Error writing to file"
ok_message:   .asciiz     "All right"
value:        .space      30

.text

write:        daddi        r14, r0, params_sys4
               sw          $s2, params_sys4(r0)
               daddi        $s1, r0, value
```

```

        sw          $s1, ind_value(r0)
        syscall     4
        daddi       $s0, r0, -1
        daddi       $a0, r0, ok_message
        bne         r1, $s0, end
        daddi       $a0, r0, error_4

end:     jal         print_string
        syscall     0

#include    print.s

```

The first 4 rows of the .data section contain the arguments of SYSCALL 4, the file descriptor of the from which we must read, the memory address from where the SYSCALL must read the bytes to write, the number of bytes to write. We give labels to those parameters that must be accessed later. Next we put, as usual, the strings containing the error message and the success message.

In the .text section, we save the params_sys4 address to register r14, we save in the memory cells for the SYSCALL parameters the file descriptor (that we suppose to have in \$s2) and the address from where we must take the bytes to write.

Next we can call SYSCALL 3, and then we call the print_string function passing as argument error_3 or ok_message, according to the success of the operation.

SYSCALL 5

Example program that contains a function that prints to standard output the string contained in \$a0:

```

        .data
params_sys5: .space 8

        .text
print_string:
        sw          $a0, params_sys5(r0)
        daddi       r14, r0, params_sys5
        syscall     5
        jr          r31

```

The second row is used to save space for the string that must be printed by the SYSCALL, that is filled by the first instruction of the .text section, that assumes that in \$a0 there's the address of the string to be printed.

The next instruction puts in r14 the address of this string, and then we can call SYSCALL 5 and print the string. The last instruction sets the program counter to the content of r31, as the usual MIPS calling convention states.

A more complex usage example of SYSCALL 5

SYSCALL 5 uses a not-so-simple arguments passing mechanism, that will be shown in the following example:

```

        .data
format_str: .asciiz  "%dth of %s:\n%s version %i.%i is being tested!"
s1:         .asciiz  "June"
s2:         .asciiz  "EduMIPS64"
fs_addr:    .space   4
            .word    5
s1_addr:    .space   4
s2_addr:    .space   4
            .word    0
            .word    5

        .code
test:
        daddi       r5, r0, format_str
        sw          r5, fs_addr(r0)
        daddi       r2, r0, s1

```

```
daddi    r3, r0, s2
sd       r2, s1_addr(r0)
sd       r3, s2_addr(r0)
daddi    r14, r0, fs_addr
syscall  5
syscall  0
```

The address of the format string is put into R5, whose content is then saved to memory at address `fs_addr`. The string parameters' addresses are saved into `s1_addr` and `s2_addr`. Those two string parameters are the ones that match the two `%s` placeholders in the format string.

Looking at the memory, it's obvious that the parameters matching the placeholders are stored immediately after the address of the format string: numbers match integer parameters, while addresses match string parameters. In the `s1_addr` and `s2_addr` locations there are the addresses of the two strings that we want to print instead of the `%s` placeholders.

The execution of the example will show how `SYSCALL 5` can handle complex format strings like the one stored at `format_str`.